

Homework 2

by Franklin Marsh

MATH0154

Computational Statistics

with Prof. 🥰Gabe Chandler🥰

Problem 1

Part (a)

$$P(k + 1) = \frac{n!}{(k+1)!(n-k-1)!} \times p^{k+1}(1-p)^{n-k-1} = \frac{n!}{k!(n-k)!} \times p^k(1-p)^{n-k} \times c$$

If we solve this equation for c and substitute, here is what we get:

$$P(k + 1) = \frac{p \times (n-k)}{(k+1)(1-p)} \times P(k)$$

We will write this into our NextTerm function.

Part (b)

```
In [469]: NextTerm <- function(prev,p,n,k) {  
  
  # Calculates binomial probability  $p(n,k+1)$  if given  $p(n,k)$   
  #  
  # Args:  
  #   prev:  $p(n,k)$   
  #   p: probability between 0 and 1, 0.5 for coin flip.  
  #   n: total number of trials  
  #   k: number of succesful trials  
  # Returns:  
  #    $p(n,k+1)$ : next binomial probability  
  
  return((p*(n-k))/((k+1)*(1-p)))*prev) #return the formula  
}
```

This next function will draw m random variables from the binomial distribution using the discrete inverse transform.

```

In [517]: DrawBinomial <- function(p,m,n) {

    #Performs m experiments on a binomial distribution with n trials a
nd probability p
    #
    # Args:
    #   p: probability between 0 and 1, 0.5 for coin flip.
    #   n: total number of trials per experiment
    #   m: total number of experiments
    #
    # Returns:
    #   out: results of experiments in a vector of length m.

    out <- c() #initialize the output vector as empty

    for (num in (1:m)) { #for each trial

        P <- c((1-p)^(n)) #calculate the starting P0
        U <- runif(1) #draw a uniform
        CDF <- P #add P0 to the cdf
        k <- 0 #start k at 0

        while (U > CDF) { #while the uniform is still greater than the
accumulated CDF
            P <- NextTerm(P, p, n, k) #calculate the next probability
            CDF <- P + CDF #and add it to the cdf
            k <- k + 1 # add 1 to k
        }

        out[num] <- k #append k to the output vector
    }

    return(out)
}

```

Draw some random binomials! Looks good! 100

```
In [523]: hist(DrawBinomial(0.5,100,10))
```



Problem 2

Part (a)

The CDF of a single exponential distribution is $F(x) = 1 - e^{-\lambda x}$.

Using the symmetry of the uniform distribution, $U = 1 - e^{-\lambda x} \Rightarrow x = \frac{\log(u)}{\lambda}$

```
In [376]: DoubleExponential <- function(lambda, n_iters) {

  # Draw random variables from a double exponential distribution with
  # parameter lambda
  #
  # Args:
  #   lambda: double exponential scale factor.
  #   n_iters: number of desired random variates.
  #
  # Returns:
  #   x: vector of random variates with length m.

  x <- c()

  sign <- runif(n_iters) #initialize a random uniform to choose sign
  val <- runif(n_iters) #initialize a random uniform to generate value
  ue


  for (i in (1:n_iters)) {

    if (sign[i] > 0.5) {
      x[i] <- log(val[i])/lambda
    }
    else if (sign[i] < 0.5) {
      x[i] <- -log(val[i])/lambda
    }

  }

  return (x)
}
```

```
In [198]: hist(DoubleExponential(3, 1000), breaks = 100)
```




Looks good!

Part (b)


We'll need the pdf of both the masking distribution (Double Exponential Distribution), and the target distribution (Normal Distribution).

```
In [210]: plot(seq(-3,3,0.05),dnorm(seq(-3,3,0.05)))
```



```
In [524]: DE_pdf <- function(x, lambda = 1) {  
  
    # Calculates Double Exponential probability density function  
    #  
    # Args:  
    #     x: evaluate the pdf for this x-value.  
    #     lambda: the scale factor of the double exponential distribut  
    ion.  
  
    return((lambda/2.0)*exp(1)^(-lambda*abs(x)))  
  
}
```

```
In [217]: plot(DE_pdf(seq(-3,3,0.05),4))
```

```
In [225]: FindC <- function(maskpdf, targetpdf, param1, prange = seq(-3,3,0.01),
by = 0.01) {

    # Finds the normalizing coefficient c for drawing using the maskin
g technique
    #
    # Args:
    #     maskpdf:
    #     targetpdf:
    #     param1:
    #     prange:
    #     by:
    # Returns:
    #     c: normalizing coefficient

    grid <- seq(min(prange),max(prange), by)


    c <- max((targetpdf(grid))/(maskpdf(grid, param1)))

    return(c)
}
```

```
In [240]: grid <- seq(0.5,1.5, 0.005) #test over a range of many lambda values
Cgrid <- c() #create empty array to store c calculate for many lambda
values
i <- 0 # initialize empty variable to populate Cgrid

for (lambda in grid) {
    i <- i + 1
    Cgrid[i] <- FindC(DE_pdf, dnorm, lambda)
}


plot(grid,Cgrid, log = 'y')
```



This grid search tells us that the C to use is about 1.3154... We also seem to have to set $\lambda = 1$ to achieve this low C .

```
In [298]: min(Cgrid) #find the minimum of this grid search. This is our C  
1.31548924695891
```

```
In [459]: abciissa <- seq(-3,3,0.05) #create some abciissa values to plot two pdfs  
.  
  
plot(abciissa,dnorm(abciissa), type = 'l')  
par(new = T)  
lines(abciissa,DE_pdf(abciissa, 1.0), type = 'l')  
par(new = F)
```



Above is are the two PDFs drawn on the same axis.

Part (c)

```
In [293]: MaskingDraw <- function(maskdraw, maskpdf, goalpdf, c, n_iters) {

    # Draw from pdf (goalpdf) using the masking technique
    #
    # Args:
    #   maskdraw: a function that draws from the masking dist.
    #   maskpdf: a function that returns the pdf of the masking dist
    #   when given an x.
    #   goalpdf: the pdf of the function that we are trying draw from
    #   c: the c normalizing constant for maskpdf and goalpdf
    #   n_iters: number of points that we want to draw
    #
    # Returns:
    #   out: random numbers like they were drawn from goalpdf

    out <- c() #initialize empty output

    for (i in (1:n_iters)) {

        pull <- maskdraw(lambda = 1, 1) #the initial value we are pulling from the masking distribution
        crit <- (goalpdf(pull))/(c*maskpdf(pull))

        U <- runif(1)

        if (U <= crit) {
            out[i] <- pull
        }
    }
    return(out)
}
```


```
In [295]: sd(na.omit(MaskingDraw(DoubleExponential, DE_pdf, dnorm, min(Cgrid), 10000)))
```

1.00031977983949

The resulting distribution has a standard deviation of one. Here is how the histogram looks:

```
In [299]: hist(na.omit(MaskingDraw(DoubleExponential, DE_pdf, dnorm, min(Cgrid),  
10000)))
```

ram of na.omit(MaskingDraw(DoubleExponential, DE_pdf, dnorm, min(Cgr



na.omit(MaskingDraw(DoubleExponential, DE_pdf, dnorm, min(Cgrid), 10000))

```
In [300]: mask_draws <- MaskingDraw(DoubleExponential, DE_pdf, dnorm, min(Cgrid),  
10000)
```

```
In [303]: rejection_probability <- (length(mask_draws) - length(na.omit(mask_draws))/  
length(mask_draws))  
print(rejection_probability)
```

```
[1] 0.2332
```

For $C = 1.3155$, about 23.32% of the draws were rejected by the masking process.

Problem 3

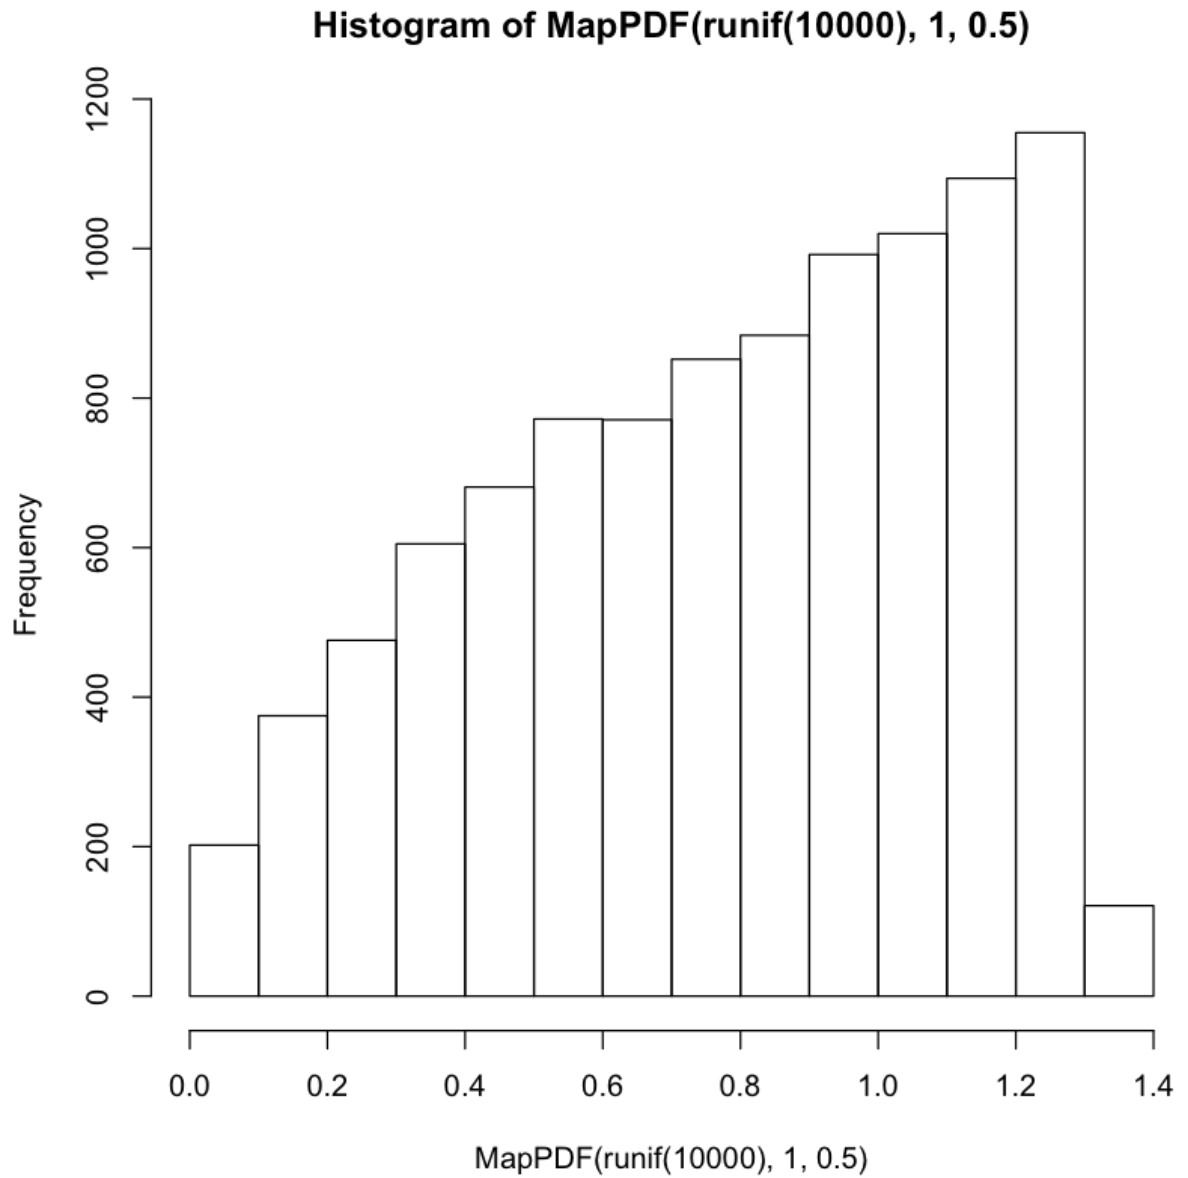
```
In [483]: MapPDF <- function(u, alpha, theta) {

  # Maps uniforms to a new pdf as described in problem 3:
  #
  #  $f(x) = \alpha x^{(\theta)}$ 
  #
  # Args:
  #   u: The input uniform that is to be mapped
  #   alpha: the 'alpha' parameter in the pdf definition
  #   theta: the 'theta' parameter in the pdf definition
  # Returns:
  #   out: a float of the mapped uniform

  return(((u*(theta + 1))/alpha)^(1/(theta + 1)))
}
```

We are taking the function $x = \left(\frac{u(\theta+1)}{\alpha}\right)^{\frac{1}{\theta+1}}$

```
In [484]: hist(MapPDF(runif(10000), 1, 0.5))
```



Problem 4


```

In [335]: pdf <- function(x, theta = 0, c = 2*pi) {

  # Evaluates the PDF in problem 4
  #
  # Args:
  #   x: point to evaluate pdf
  #   theta: phase offset in input function
  #   c: normalization constant, defaults to 1
  #
  # Returns:
  #   pdf(x): pdf evaluated at x

  return((1 - cos(x - theta))/(c))

}

StochasticIntegral <- function(pdf, prange = c(0,2*pi), n_iters) {

  # Performs stochastic (monte carlo) integration
  #
  # Args:
  #   pdf: the function to integrate
  #   prange: the range of x values to integrate over
  #   n_iters: number of draws to perform to evaluate the integral
  #
  # Returns:
  #   integral: the stochastic estimate of the integral

  V <- max(prange) - min(prange)
  hold <- c() #hold
  i <- 0
  for (U in runif(n_iters)*V) {
    i <- i + 1
    hold[i] <- pdf(U, 0)
  }

  return((sum(hold)/n_iters)*V)
}

```

Part (a)

```

In [179]: StochasticIntegral(pdf, n_iters = 10000)

```

6.26603711856409

Looks like the normalization constant needs to be 2π

Part (b)

The likelihood: $L(\theta) = \prod_{i=1}^n \left(\frac{1 - \cos(x_i - \theta)}{2\pi} \right)$

The log-likelihood: $l(\theta) = \log(L(\theta)) = \sum_{i=1}^n \log \left(\frac{1 - \cos(x_i - \theta)}{2\pi} \right)$

Part (c)

$$\frac{d}{d\theta} \left(\log(L(\theta)) \right) = \sum_{i=1}^n \cot \left(\frac{(x_i - \theta)}{2} \right)$$

$$\frac{d}{d^2\theta} \left(\log(L(\theta)) \right) = -\frac{1}{2} \sum_{i=1}^n \csc^2 \left(\frac{x_i - \theta}{2} \right)$$

Part (d)

```
In [346]: FindC_2 <- function(maskpdf, goalpdf, prange = c(0.01, 2*pi - 0.01), by
= 0.01) {

  # See FindC.

  grid <- seq(min(prange), max(prange), by)

  c <- max((goalpdf(grid))/(maskpdf(grid, min = min(prange), max = m
ax(prange))))

  return(c)
}
```

```
In [340]: C2 <- FindC_2(dunif, pdf, by = 0.001)
print(C2)

[1] 1.993634
```

It appears that 2 is a good choice for c.

```
In [358]: DrawPDF <- function(n_iter) {

    # A sloppy function that only draws from the pdf in this function.
Do not port to another R script.
    #
    # Args:
    #     n_iters: number of random variates desired
    # Returns:
    #     out: vector of length n_iters of the random variates.

    out <- c()


    for (i in (1:n_iter)) {

        pull <- runif(1,min = 0, max = 2*pi) #the initial value we are
pulling from the masking distribution
        crit <- (2*pi*pdf(pull))/C2

        U <- runif(1)

        if (U <= crit) {
            out[i] <- pull
        }
    }
    return(out)
}
```

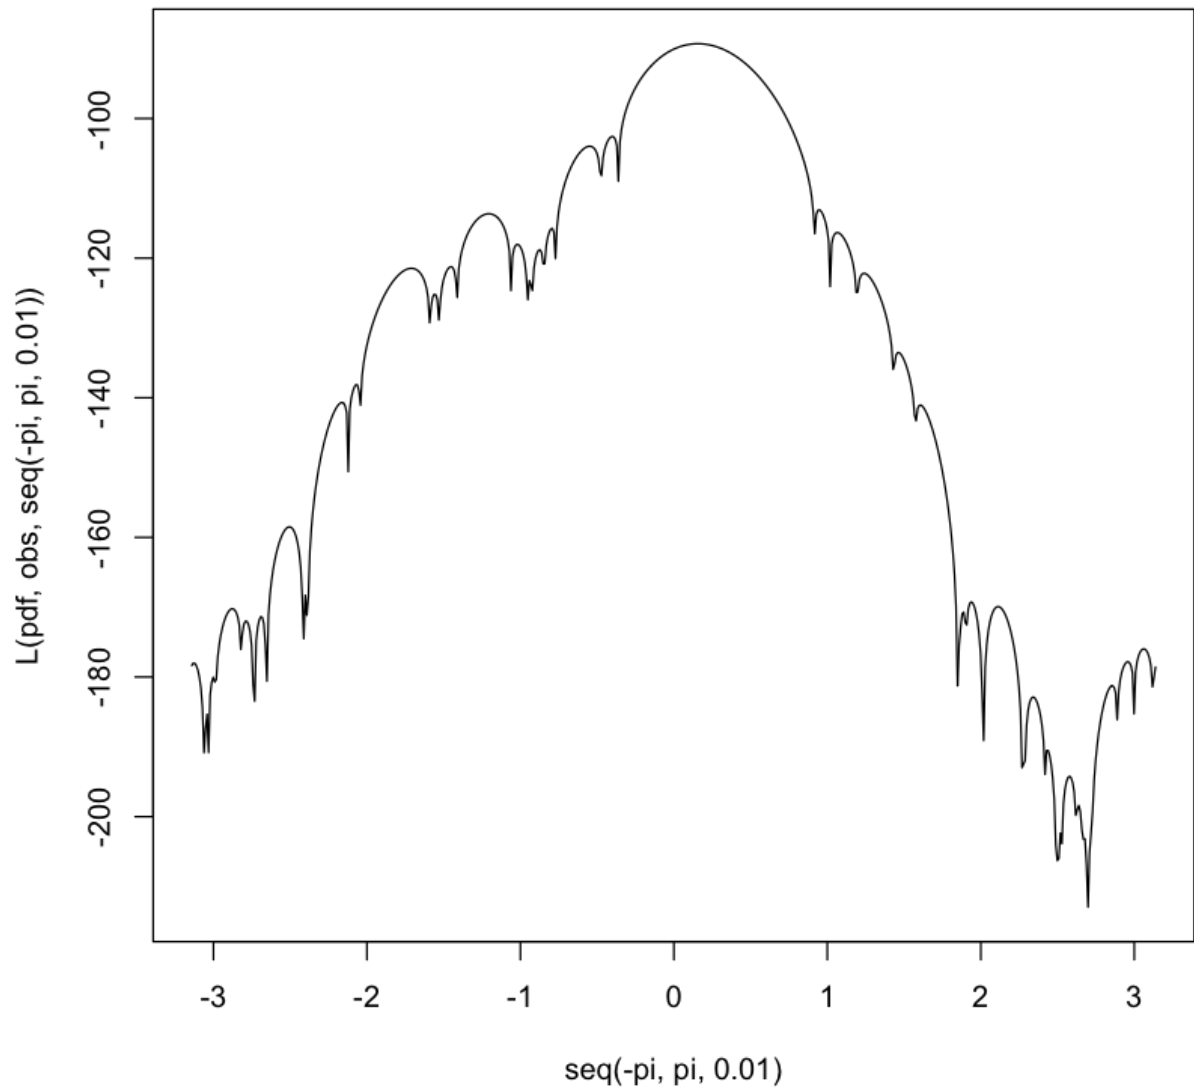
```
In [360]: obs <- DrawPDF(100)  
hist(obs)
```



Above is a histogram of the observations.

```
In [427]: L <- function(pdf, obs, trange, ll = TRUE) {  
  
  # Computes log likelihood as a function of theta (trange)  
  #  
  # Args:  
  #   pdf: input pdf  
  #   obs: observations to be used in calculating pdf  
  #   trange: range of theta values to plot over  
  # Returns:  
  #   out: range of theta  
  #  
  
  out <- c() #initialize output  
  i <- 0 #initialize counter for list index  
  
  for (t in trange) { #for every theta  
    i <- i + 1  
    if (ll == TRUE) { # if you want the log-likelihood  
      out[i] <- log(prod(pdf(obs, theta = t), na.rm = TRUE))  
    }  
    else if (ll == FALSE) {  
      out[i] <- prod(pdf(obs, theta = t), na.rm = TRUE)  
    }  
  }  
  
  return(out)  
}
```

```
In [525]: plot(seq(-pi, pi, 0.01),L(pdf,obs, seq(-pi, pi, 0.01)), type = 'l')
```



This function appears to be at maximum when $\theta = 0$, which makes sense because we used $\theta = 0$ to generate the data!

Part (e)

The iterative process in Newton's method is :

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

So, we need to be able to evaluate the function's 1st derivative and its 2nd derivative at every point.

Thankfully we already have these:

$$l(\theta) = \log(L(\theta)) = \sum_{i=1}^n \log \left(\frac{1 - \cos(x_i - \theta)}{2\pi} \right)$$

$$\frac{d}{d\theta} \left(\log(L(\theta)) \right) = \sum_{i=1}^n \cot \left(\frac{(x_i - \theta)}{2} \right)$$


$$\frac{d}{d^2\theta} \left(\log(L(\theta)) \right) = -\frac{1}{2} \sum_{i=1}^n \csc^2 \left(\frac{x_i - \theta}{2} \right)$$

```
In [454]: l_theta <- function(theta, obs) {  
  
    # calculate log-likelihood of observations, given observations and  
    theta.  
  
    terms <- c()  
    i <- 0  
  
    for (ob in obs) {  
        i <- i + 1  
        terms[i] <- log((1 - cos(ob - theta))/(2*pi))  
    }  
  
    return(sum(terms))  
}  
  
dl_theta <- function(theta, obs) {  
  
    # calculate 1st derivative log-likelihood of observations, given o  
    bservations and theta.  
  
    terms <- c()  
    i <- 0  
  
    for (ob in obs) {  
        i <- i + 1  
        terms[i] <- tan((ob - theta)/(2))(-1)  
    }  
  
    return(sum(terms))  
}  
  
ddl_theta <- function(theta, obs) {  
  
    # calculate 2nd derivative log-likelihood of observations, given o  
    bservations and theta.  
  
    terms <- c()  
    i <- 0  
  
    for (ob in obs) {  
        i <- i + 1  
        terms[i] <- (-1/2)*sin((ob - theta)/(2))(-2)  
    }  
  
    return(sum(terms))  
}
```



```
In [456]: NewtonsMethod <- function(ifunc, difunc, x0, obs, tol = 1E-5, maxn = 30) {  
  
    # Perform Newton's method to find a root, given the 1st and second  
    derivatives of that function.  
  
    x <- c(x0)  
    i <- 1  
    while (i <= maxn) {  
        i <- i + 1  
        x[i] <- x[i-1] - (ifunc(x[i-1], na.omit(obs)))/(difunc(x[i-1],  
na.omit(obs)))  
        if (abs(x[i-1] - x[i]) < tol) break  
    }  
    return(x)  
}
```

```
In [463]: plot(NewtonsMethod(dl_theta, ddl_theta, 4, obs, maxn = 300 ))
```



As we can see in the plot above, the solver finds an approximate value of a maximum very quickly, however, it is not able to find the global maximum, because there are many local maxima, and the solver just finds the nearest local maximum 😞.

In []: